

```
In [1]: #Import relevant Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

#Set style
plt.style.use('seaborn-whitegrid')
```

## Load Datasets

### Dataset 1 -- King County House Prices (2014 - 2015)


#### a) Load and Preview the Data

```
In [2]: #Load dataset into DataFrame and preview
kc = pd.read_csv("Housing_Prices_Modeling/Data/kc_house_data.csv")
kc.head()
```

Out[2]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfr
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	N
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns



In [3]: *#Check DataFrame summaries*

```

kc.info()
kc.describe()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
id                21597 non-null int64
date              21597 non-null object
price             21597 non-null float64
bedrooms          21597 non-null int64
bathrooms         21597 non-null float64
sqft_living       21597 non-null int64
sqft_lot          21597 non-null int64
floors            21597 non-null float64
waterfront        19221 non-null float64
view              21534 non-null float64
condition         21597 non-null int64
grade             21597 non-null int64
sqft_above        21597 non-null int64
sqft_basement     21597 non-null object
yr_built          21597 non-null int64
yr_renovated      17755 non-null float64
zipcode           21597 non-null int64
lat               21597 non-null float64
long              21597 non-null float64
sqft_living15     21597 non-null int64
sqft_lot15        21597 non-null int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB

```

Out[3]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot
<b>count</b>	2.159700e+04	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04
<b>mean</b>	4.580474e+09	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04
<b>std</b>	2.876736e+09	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04
<b>min</b>	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02
<b>25%</b>	2.123049e+09	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03
<b>50%</b>	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
<b>75%</b>	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04
<b>max</b>	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06

## b) Clean Data

```
In [4]: #Check column names  
print(kc.columns)
```

```
Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',  
      'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',  
      'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',  
      'lat', 'long', 'sqft_living15', 'sqft_lot15'],  
      dtype='object')
```

```
In [5]: #Examine sqft_basement values  
print(kc.sqft_basement.value_counts())
```

```
0.0      12826  
?         454  
600.0     217  
500.0     209  
700.0     208  
...  
1481.0      1  
704.0      1  
172.0      1  
1852.0      1  
2500.0      1  
Name: sqft_basement, Length: 304, dtype: int64
```

```
In [6]: #Rows to drop
to_drop = kc.loc[kc['sqft_basement'] == '?'].index

#Drop rows
kc.drop(to_drop, inplace = True)

#Change the Series type to float
kc['sqft_basement'] = kc['sqft_basement'].astype(str).astype(float)

#Check DataFrame
kc.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21143 entries, 0 to 21596
Data columns (total 21 columns):
id                21143 non-null int64
date              21143 non-null object
price             21143 non-null float64
bedrooms          21143 non-null int64
bathrooms         21143 non-null float64
sqft_living       21143 non-null int64
sqft_lot          21143 non-null int64
floors            21143 non-null float64
waterfront        18804 non-null float64
view              21082 non-null float64
condition         21143 non-null int64
grade             21143 non-null int64
sqft_above        21143 non-null int64
sqft_basement     21143 non-null float64
yr_built          21143 non-null int64
yr_renovated      17389 non-null float64
zipcode           21143 non-null int64
lat               21143 non-null float64
long              21143 non-null float64
sqft_living15     21143 non-null int64
sqft_lot15        21143 non-null int64
dtypes: float64(9), int64(11), object(1)
memory usage: 3.5+ MB
```

```
In [7]: #Check for duplicate rows
duplicate_rows = kc[kc.duplicated()]
print(duplicate_rows.shape)

(0, 21)
```

```
In [8]: #Check for null values
print(kc.isna().sum())
```

```
id                0
date              0
price             0
bedrooms          0
bathrooms         0
sqft_living       0
sqft_lot          0
floors            0
waterfront       2339
view              61
condition         0
grade             0
sqft_above        0
sqft_basement     0
yr_built          0
yr_renovated      3754
zipcode           0
lat               0
long              0
sqft_living15     0
sqft_lot15        0
dtype: int64
```

```
In [9]: #Examine the values in the columns with null values
print('View of the Waterfront values: ', kc['waterfront'].value_counts())
print('Year Renovated values: ', kc['yr_renovated'].value_counts())
print('View values: ', kc['view'].value_counts())
```


```
View of the Waterfront values: 0.0    18662
1.0     142
Name: waterfront, dtype: int64
Year Renovated values: 0.0    16666
2014.0     69
2003.0     31
2013.0     31
2007.0     30
...
1953.0     1
1944.0     1
1934.0     1
1971.0     1
1959.0     1
Name: yr_renovated, Length: 69, dtype: int64
View values: 0.0    19018
2.0     930
3.0     496
1.0     327
4.0     311
Name: view, dtype: int64
```

```
In [10]: #Drop the unnecessary and/or uninteresting columns / columns with null values
kc.drop(['id', 'date', 'view', 'yr_renovated'], axis = 1, inplace = True)

#Check DataFrame
kc.head(2)
```

Out[10]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	condition	grade	sq
0	221900.0	3	1.00	1180	5650	1.0	NaN	3	7	
1	538000.0	3	2.25	2570	7242	2.0	0.0	3	7	



```
In [11]: #Drop the remaining null values by the row
kc.dropna(inplace = True)

#Re-Check for null values
print(kc.isna().sum())
```

```
price          0
bedrooms       0
bathrooms      0
sqft_living    0
sqft_lot       0
floors         0
waterfront     0
condition      0
grade          0
sqft_above     0
sqft_basement  0
yr_built       0
zipcode        0
lat            0
long           0
sqft_living15  0
sqft_lot15     0
dtype: int64
```

## Dataset 1 -- King County House Prices (2014 - 2015)

### a) Load and Preview the Data

In [12]: *#Load dataset into DataFrame and preview*

```
cpi = pd.read_csv("Housing_Prices_Modeling/Data/Common_Points_of_Interest_for_King_County___common_interest_point.csv")
cpi.head()
```

Out[12]:

	X	Y	OBJECTID	FEATURE_ID	ESITE	CODE	NAME	ABB_NAME	AC
0	-122.286944	47.499985	1	6002948	0.0	600	Green River Trail Site - Tukwila	Green River Trail Site - Tukwila	1169
1	-122.305465	47.635532	2	828	0.0	600	Interlaken Park	Interlaken Park	1169
2	-122.211064	47.405961	3	374	0.0	600	Garrison Creek Park	Garrison Creek Park	1169
3	-121.912156	47.650466	4	1891	124849.0	390	Carnation Library	Carnation Lib	1169
4	-122.295038	47.441348	5	1817	401027.0	60	Sea-Tac Office Center	Sea-Tac Office Center	1169

In [13]: *#Check the DataFrame summary*

```
cpi.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6491 entries, 0 to 6490
Data columns (total 10 columns):
X                6491 non-null float64
Y                6491 non-null float64
OBJECTID         6491 non-null int64
FEATURE_ID       6491 non-null int64
ESITE            6484 non-null float64
CODE             6491 non-null int64
NAME             6491 non-null object
ABB_NAME         6491 non-null object
ADDRESS          6489 non-null object
ZIPCODE          6491 non-null int64
dtypes: float64(3), int64(4), object(3)
memory usage: 507.2+ KB
```

## b) Clean Data

```
In [14]: #Rename all column names to lowercase
        cpi.columns = [x.lower() for x in cpi.columns]

        #Check column names
        print(cpi.columns)
```

```
Index(['x', 'y', 'objectid', 'feature_id', 'esite', 'code', 'name', 'abb_name',
       'address', 'zipcode'],
      dtype='object')
```

```
In [15]: #Remove uninteresting columns
        cpi.drop(['x', 'y', 'objectid', 'feature_id', 'esite'], axis = 1, inplace = True)
```

```
In [16]: #Check for duplicate rows in the dataset
        duplicate_rows = cpi[cpi.duplicated()]
        print(duplicate_rows.shape)

        #Checkk for null values
        print(cpi.isna().sum())
```

```
(1, 5)
code      0
name      0
abb_name  0
address   2
zipcode   0
dtype: int64
```

```
In [17]: #Drop duplicate rows
        cpi.drop_duplicates(inplace = True)

        #Drop the 2 rows with null values
        cpi.dropna(inplace = True)
```

```
In [18]: #Re-Check the values
        print(cpi.isna().sum())
```

```
code      0
name      0
abb_name  0
address   0
zipcode   0
dtype: int64
```



```
In [19]: #Check Zipcode values
zipcode = list(cpi.zipcode.value_counts().index)
print(sorted(zipcode))
```

```
[0, 9827, 91855, 98001, 98002, 98003, 98004, 98005, 98006, 98007, 98008, 9800
9, 98010, 98011, 98012, 98014, 98015, 98019, 98020, 98021, 98022, 98023, 9802
4, 98027, 98028, 98029, 98030, 98031, 98032, 98033, 98034, 98035, 98036, 9803
8, 98039, 98040, 98042, 98043, 98045, 98047, 98050, 98051, 98052, 98053, 9805
4, 98055, 98056, 98057, 98058, 98059, 98063, 98065, 98068, 98070, 98072, 9807
4, 98075, 98077, 98083, 98092, 98101, 98102, 98103, 98104, 98105, 98106, 9810
7, 98108, 98109, 98110, 98112, 98115, 98116, 98117, 98118, 98119, 98121, 9812
2, 98124, 98125, 98126, 98127, 98133, 98134, 98136, 98138, 98144, 98145, 9814
6, 98148, 98154, 98155, 98158, 98160, 98164, 98165, 98166, 98168, 98174, 9817
7, 98178, 98185, 98187, 98188, 98191, 98195, 98198, 98199, 98203, 98204, 9822
4, 98288, 98366, 98391]
```

```
In [20]: #Select rows to drop
rows_to_drop = cpi.loc[(cpi['zipcode'] == 0) | (cpi['zipcode'] == 9827)].index

#Drop the rows
cpi.drop(rows_to_drop, inplace = True)
```

```
In [21]: #Check DataFrame information
cpi.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 6178 entries, 0 to 6490
Data columns (total 5 columns):
code          6178 non-null int64
name          6178 non-null object
abb_name      6178 non-null object
address       6178 non-null object
zipcode       6178 non-null int64
dtypes: int64(2), object(3)
memory usage: 289.6+ KB
```

### c) Webscrape Dataset metadata

The 'code' Series in the Common Points of Interest (cpi) DataFrame corresponds to points of interest along multiple different domain classes (e.g. Schools or Parks and Recreation).

Check Metadata:

<https://www.arcgis.com/sharing/rest/content/items/6f6dfde35681494f92da924faf7ee47c/info/metadata/metadata.xrformat=default&output=html>  
<https://www.arcgis.com/sharing/rest/content/items/6f6dfde35681494f92da924faf7ee47c/info/metadata/metadata.xrformat=default&output=html>

```
In [22]: import requests
from bs4 import BeautifulSoup
```

```
In [23]: html_page = requests.get('https://www.arcgis.com/sharing/rest/content/items/6f6dfde35681494f92da924faf7ee47c/info/metadata/metadata.xml?format=default&output=html')
soup = BeautifulSoup(html_page.content, 'html.parser')
container = soup.find('body', class_ = 'bodyText')
```

```
In [24]: #Get the text from the html source
page_text = container.get_text()

#Clean the text
#Split into a list
page_text = page_text.split('\n')

#Remove empty elements in the list
for item in page_text:
    if item == '':
        page_text.remove(item)
```

```
In [25]: #Create an empty list to append the relevant reference information on the CODE
Series in the DataFrame
temp_list = []

#Iterate over the page_text list to find the necessary information and append
to temp_list
for item in page_text:
    if 'Enumerated Domain Value: ' in item and item not in temp_list:
        temp_list.append(item)
    if 'Enumerated Domain Value Definition: ' in item and item not in temp_list:
        temp_list.append(item)

#Clean the entries in the temp_list
#Remove duplicates & create final lists of keys and values to create a reference
dictionary
#Create two empty keys and values lists
keys = []
values = []

#Assemble the keys and values list
for item in temp_list:
    if temp_list.index(item) % 2 == 0:
        item = item.replace('Enumerated Domain Value:', '')
        keys.append(int(item))
    else:
        item = item.replace('Enumerated Domain Value Definition: ', '')
        values.append(item)

#Check the lengths of both lists match
print(len(keys), len(values))

# Create and assemble the reference dictionary
ref = {}
for index in range(0, len(keys)):
    ref[keys[index]] = values[index]

#Check the new reference dictionary
ref
```

49 49

```

Out[25]: {30: 'Airport',
        60: 'Non-Government Building',
        61: 'Building - City Government',
        62: 'Building - County Government',
        63: 'Building - State Government',
        64: 'Building - Federal Government',
        65: 'Police Station',
        66: 'Fire Station',
        67: 'Medic Units',
        90: 'Community area, Business Center or neighborhood',
        120: 'Cemetery',
        150: 'Chamber of Commerce',
        180: 'City Hall',
        210: 'Department of Motor Vehicles',
        240: 'Entertainment and Sport facility',
        270: 'Fairground',
        300: 'Golf Course',
        330: 'Hospital or Medical Center',
        340: 'Public Health Clinic',
        360: 'Hotel or Motel',
        390: 'Library',
        420: 'Major employment center or large business',
        450: 'Military installation',
        480: 'Museum',
        500: 'Transit center',
        510: 'Other transportation center',
        520: 'Fare outlet - All type',
        530: 'Fare outlet - Limited type',
        540: 'Park and Pool',
        570: 'Park and Ride',
        580: 'Bike Lockers',
        581: 'Electrical outlets',
        600: 'Parks and Recreation',
        630: 'Pier or Terminal',
        660: 'School - Elementary',
        661: 'School - Junior High or Middle',
        662: 'School - High',
        663: 'School - College or University',
        664: 'School - Alternative',
        665: 'School - Other facility',
        666: 'Schools - K thru 12',
        690: 'Shopping center',
        720: 'Winery',
        902: 'Bike shop',
        903: 'Farmers Market',
        904: 'Public Access farm',
        999: 'General reference feature',
        350: 'Food Facility',
        68: 'Public Safety Answering Point (PSAP)'}

```

```

In [26]: #Bin the reference points of interest into categories with the list of corresponding codes
domains_of_interest = {'public_safety' : [65, 66, 67, 68],
                        'government_building' : [61, 62, 63, 64, 150, 180, 450],
                        'schools' : [660, 661, 662, 663, 664, 665, 666],
                        'parks_and_recreation' : [540, 570, 600],
                        'food_and_restaurants' : [720, 903, 350],
                        'shopping_and_entertainment' : [90, 240, 270, 300, 390,
480, 690, 902],
                        'hubs_of_transport' : [30, 500, 510, 520, 530, 630],
                        'medical' : [67, 330, 340],
                        'other' : [60, 120, 210, 360, 580, 581, 630, 904, 994]}

In [27]: zipcodes = sorted(list(cpi.zipcode.value_counts().index))
cpi_per_area = {}

for zipcode in zipcodes:
    #Create count dictionary
    temp_dict = {'public_safety' : 0,
                  'government_building' : 0,
                  'schools' : 0,
                  'parks_and_recreation' : 0,
                  'food_and_restaurants' : 0,
                  'shopping_and_entertainment' : 0,
                  'hubs_of_transport' : 0,
                  'medical' : 0,
                  'other' : 0}

    #Create list of codes occurring in the zipcode
    codes_list = list(cpi.loc[cpi['zipcode'] == zipcode, 'code'].values)

    #Iterate through codes and the domain_of_interest reference list to count the number of points of interest per zipcode
    for code in codes_list:
        for domain in domains_of_interest:
            if code in domains_of_interest[domain]:
                temp_dict[domain] += 1

    #Append the zipcode specific dictionary to the overall dict
    cpi_per_area[zipcode] = temp_dict

In [28]: #Create an empty (full zeros) DataFrame to concat with kc
a = np.zeros(shape =(kc.shape[0],9))
tempdf = pd.DataFrame(a, columns = list(domains_of_interest.keys()))

In [29]: #Reset the index of the 1st DataFrame kc
kc.reset_index(inplace = True)
kc.drop('index', axis = 1, inplace = True)


```

```
In [30]: #Create new, final DataFrame with the new Series of interest
kcdf = pd.concat([kc, tempdf], axis = 1)
kcdf.head()
```

Out[30]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	condition	grade	s
0	538000.0	3	2.25	2570	7242	2.0	0.0	3	7	
1	180000.0	2	1.00	770	10000	1.0	0.0	3	6	
2	604000.0	4	3.00	1960	5000	1.0	0.0	5	7	
3	510000.0	3	2.00	1680	8080	1.0	0.0	3	8	
4	1230000.0	4	4.50	5420	101930	1.0	0.0	3	11	

5 rows × 26 columns



```
In [31]: #Populate the new Series in the King County DataFrame with the reference zipdf
for index in range(len(kcdf)):
```

```
    #Check if the zipcode is present in the cpi_per_area reference
    if kcdf.zipcode[index] in cpi_per_area:
```

```
        #Access the reference dictionary for the zipcode
        reference_dict = cpi_per_area[kcdf.zipcode[index]]
```

```
        #Change the value to 1 if the cpi is present, 0 if not present in the
        area
```

```
        for key in reference_dict.keys():
            if reference_dict[key] == 0:
                kcdf.loc[index, key] = 0
            else:
                kcdf.loc[index, key] = 1
```

```
    #If not, replace with a null value
```

```
    else:
        for key in reference_dict.keys():
            kcdf.loc[index, key] = np.nan
```

```
In [32]: #Check for null values
kcdf.isna().sum()
```

```
Out[32]: price                0
bedrooms                    0
bathrooms                   0
sqft_living                 0
sqft_lot                    0
floors                      0
waterfront                  0
condition                   0
grade                       0
sqft_above                  0
sqft_basement               0
yr_built                    0
zipcode                     0
lat                          0
long                        0
sqft_living15               0
sqft_lot15                  0
public_safety               0
government_building         0
schools                     0
parks_and_recreation        0
food_and_restaurants         0
shopping_and_entertainment   0
hubs_of_transport            0
medical                     0
other                       0
dtype: int64
```


## Explore Finalized DataFrame

```
In [33]: #Preview the DataFrame
kcdf.head()
```

```
Out[33]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	condition	grade	s
0	538000.0	3	2.25	2570	7242	2.0	0.0	3	7	
1	180000.0	2	1.00	770	10000	1.0	0.0	3	6	
2	604000.0	4	3.00	1960	5000	1.0	0.0	5	7	
3	510000.0	3	2.00	1680	8080	1.0	0.0	3	8	
4	1230000.0	4	4.50	5420	101930	1.0	0.0	3	11	

5 rows × 26 columns



In [34]: *#Check the Summaries*


```
kcdf.info()
kcdf.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18804 entries, 0 to 18803
Data columns (total 26 columns):
price                18804 non-null float64
bedrooms             18804 non-null int64
bathrooms            18804 non-null float64
sqft_living          18804 non-null int64
sqft_lot             18804 non-null int64
floors               18804 non-null float64
waterfront           18804 non-null float64
condition            18804 non-null int64
grade                18804 non-null int64
sqft_above           18804 non-null int64
sqft_basement        18804 non-null float64
yr_built             18804 non-null int64
zipcode              18804 non-null int64
lat                  18804 non-null float64
long                 18804 non-null float64
sqft_living15        18804 non-null int64
sqft_lot15           18804 non-null int64
public_safety        18804 non-null float64
government_building  18804 non-null float64
schools              18804 non-null float64
parks_and_recreation 18804 non-null float64
food_and_restaurants 18804 non-null float64
shopping_and_entertainment 18804 non-null float64
hubs_of_transport    18804 non-null float64
medical              18804 non-null float64
other                18804 non-null float64
dtypes: float64(16), int64(10)
memory usage: 3.7 MB
```

Out[34]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	
<b>count</b>	1.880400e+04	18804.000000	18804.000000	18804.000000	1.880400e+04	18804.000000	18804
<b>mean</b>	5.418399e+05	3.374388	2.117541	2083.155499	1.509805e+04	1.494522	
<b>std</b>	3.730331e+05	0.927297	0.769623	923.070881	4.102504e+04	0.539777	
<b>min</b>	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000	
<b>25%</b>	3.215000e+05	3.000000	1.750000	1430.000000	5.048000e+03	1.000000	
<b>50%</b>	4.500000e+05	3.000000	2.250000	1920.000000	7.629500e+03	1.500000	
<b>75%</b>	6.436125e+05	4.000000	2.500000	2550.000000	1.072075e+04	2.000000	
<b>max</b>	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	

8 rows × 26 columns





```
In [35]: continuous = ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'sqft_above', 'sqft_basement', 'lat', 'long', 'sqft_living15', 'sqft_lot15']
categorical = ['floors', 'waterfront', 'condition', 'grade', 'yr_built', 'public_safety', 'government_building', 'food_and_restaurants', 'shopping_and_entertainment', 'hubs_of_transport', 'medical', 'other']
```

```
In [36]: #Normalize/standardize the data
for column in continuous:
    if column != 'price':
        kcdf[column] = (kcdf[column] - kcdf[column].mean()) / kcdf[column].std()
()
```

## Exploratory Data Analysis & Visualization

```
In [37]: kcdf.hist(figsize = (20, 20), color = 'mediumaquamarine', edgecolor = 'black'
);
```




## Zooming in on House Sale Prices

```
In [38]: #Visualize house sales in King County in the dataset
plt.figure(figsize = (15, 6))
kc.price.plot.hist(bins = 100, color = 'c', edgecolor = 'black')

#Format the x and y axis
plt.xticks(fontsize = 12, rotation = -30)
plt.yticks(fontsize = 12)

#Format the titles
plt.title('House Sale Prices in King County (2014/15) \n', fontsize = 20, font
weight = 'bold')
plt.xlabel('\n Price', fontsize = 16, fontweight = 'bold')
plt.xticks()
plt.ylabel('Number of Houses \n', fontsize = 16, fontweight = 'bold')

plt.show();
```




```
In [82]: #Looking at the Zipcode effect on price
plt.figure(figsize = (25, 8))

#Create plot
kc.groupby('zipcode').mean()['price'].sort_values().plot(kind = 'bar',
                                                         color = 'orange',
                                                         edgecolor = 'black')

#Format plot
plt.title('Comparison of Mean Sale Price per Zipcode \n', fontweight = 'bold',
          fontsize = 20)
plt.xlabel('\n Zipcode', fontsize = 16, fontweight = 'bold')
plt.xticks(rotation = 90, fontsize = 12)
plt.ylabel('Average Price Per House \n', fontsize = 16, fontweight = 'bold')
plt.yticks(fontsize = 12)

plt.show();
```




```
In [41]: #Looking at the Sales numbers by zipcode
plt.figure(figsize = (20, 8))

#Create plot
kc.zipcode.value_counts().sort_values().plot(kind = 'bar', color = 'orange', edgecolor = 'black')

#Format plot
plt.title('Comparison of Number of House Sales per Zipcode \n', fontweight = 'bold', fontsize = 20)
plt.xlabel('\n Zipcode', fontsize = 16, fontweight = 'bold')
plt.xticks(rotation = 90, fontsize = 12)
plt.ylabel('Number of Houses \n', fontsize = 16, fontweight = 'bold')
plt.yticks(fontsize = 12)

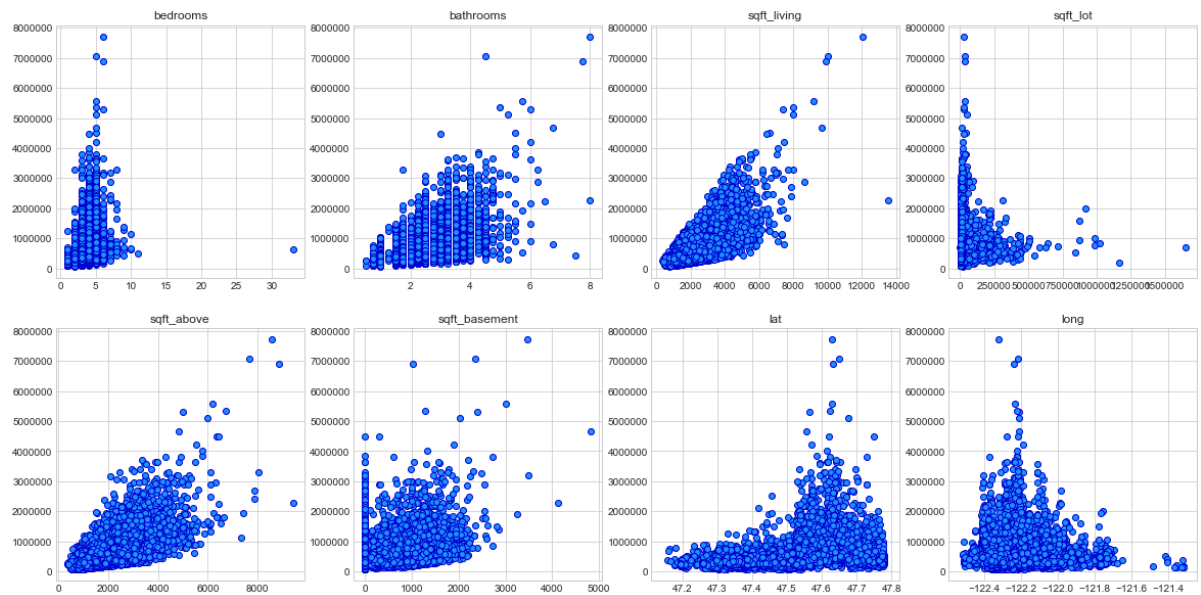
plt.show();
```



```
In [42]: #Visualize the relationships of continuous variables vis a vis price
fig, axes = plt.subplots(nrows = 2, ncols = 4, figsize = (20, 10))

#List of variables to visualize
variables = continuous
variables.remove('price')

for col, ax in zip(continuous, axes.flatten()):
    ax.scatter(kc[col], kc.price, color = 'dodgerblue', edgecolor = 'mediumblue')
    ax.set_title(col)
```



## Prepare Data for Modeling

```
In [43]: kcdf.nunique()
```

```
Out[43]: price                3354  
bedrooms                    12  
bathrooms                   29  
sqft_living                 983  
sqft_lot                   8935  
floors                      6  
waterfront                  2  
condition                   5  
grade                      11  
sqft_above                 896  
sqft_basement              299  
yr_built                   116  
zipcode                    70  
lat                        4924  
long                       739  
sqft_living15              744  
sqft_lot15                 7976  
public_safety              2  
government_building        2  
schools                    1  
parks_and_recreation       1  
food_and_restaurants       2  
shopping_and_entertainment 2  
hubs_of_transport          2  
medical                    2  
other                      2  
dtype: int64
```

```
In [44]: #Drop unsuitable columns  
kcdf.drop(['yr_built', 'zipcode', 'schools', 'parks_and_recreation'], axis = 1  
, inplace = True)  
categorical.remove('yr_built')
```

```
In [45]: #Transform categorical variables  
for col in categorical:  
    kcdf[col] = kcdf[col].astype('category')  
    kcdf[col].cat.codes
```

```
In [46]: #Create dummy variables for categorical data  
dummy = pd.get_dummies(kcdf[categorical], drop_first = True)
```

```
In [47]: #Remove original columns from dataset and add the dummy columns
kcdf.drop(categorical, axis = 1, inplace = True)

kcdf = pd.concat([kcdf, dummy], axis = 1)
kcdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18804 entries, 0 to 18803
Data columns (total 38 columns):
price                                18804 non-null float64
bedrooms                            18804 non-null float64
bathrooms                           18804 non-null float64
sqft_living                          18804 non-null float64
sqft_lot                             18804 non-null float64
sqft_above                          18804 non-null float64
sqft_basement                        18804 non-null float64
lat                                  18804 non-null float64
long                                 18804 non-null float64
sqft_living15                        18804 non-null float64
sqft_lot15                           18804 non-null float64
floors_1.5                           18804 non-null uint8
floors_2.0                           18804 non-null uint8
floors_2.5                           18804 non-null uint8
floors_3.0                           18804 non-null uint8
floors_3.5                           18804 non-null uint8
waterfront_1.0                       18804 non-null uint8
condition_2                          18804 non-null uint8
condition_3                          18804 non-null uint8
condition_4                          18804 non-null uint8
condition_5                          18804 non-null uint8
grade_4                              18804 non-null uint8
grade_5                              18804 non-null uint8
grade_6                              18804 non-null uint8
grade_7                              18804 non-null uint8
grade_8                              18804 non-null uint8
grade_9                              18804 non-null uint8
grade_10                             18804 non-null uint8
grade_11                             18804 non-null uint8
grade_12                             18804 non-null uint8
grade_13                             18804 non-null uint8
public_safety_1.0                    18804 non-null uint8
government_building_1.0              18804 non-null uint8
food_and_restaurants_1.0              18804 non-null uint8
shopping_and_entertainment_1.0        18804 non-null uint8
hubs_of_transport_1.0                 18804 non-null uint8
medical_1.0                           18804 non-null uint8
other_1.0                             18804 non-null uint8
dtypes: float64(11), uint8(27)
memory usage: 2.1 MB
```



```
In [48]: #Rename Columns
kcdf.columns = kcdf.columns.str.strip().str.replace('_', '').str.replace('0',
 '').str.replace('.', '')

kcdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18804 entries, 0 to 18803
Data columns (total 38 columns):
price                                18804 non-null float64
bedrooms                            18804 non-null float64
bathrooms                           18804 non-null float64
sqftliving                           18804 non-null float64
sqftlot                              18804 non-null float64
sqftabove                            18804 non-null float64
sqftbasement                         18804 non-null float64
lat                                  18804 non-null float64
long                                 18804 non-null float64
sqftliving15                         18804 non-null float64
sqftlot15                            18804 non-null float64
floors15                             18804 non-null uint8
floors2                              18804 non-null uint8
floors25                             18804 non-null uint8
floors3                              18804 non-null uint8
floors35                             18804 non-null uint8
waterfront1                          18804 non-null uint8
condition2                           18804 non-null uint8
condition3                           18804 non-null uint8
condition4                           18804 non-null uint8
condition5                           18804 non-null uint8
grade4                               18804 non-null uint8
grade5                               18804 non-null uint8
grade6                               18804 non-null uint8
grade7                               18804 non-null uint8
grade8                               18804 non-null uint8
grade9                               18804 non-null uint8
grade1                              18804 non-null uint8
grade11                             18804 non-null uint8
grade12                             18804 non-null uint8
grade13                             18804 non-null uint8
publicsafety1                       18804 non-null uint8
governmentbuilding1                 18804 non-null uint8
foodandrestaurants1                 18804 non-null uint8
shoppingandentertainment1           18804 non-null uint8
hubsofttransport1                   18804 non-null uint8
medical1                             18804 non-null uint8
other1                              18804 non-null uint8
dtypes: float64(11), uint8(27)
memory usage: 2.1 MB
```

## Modeling

```
In [49]: from statsmodels.formula.api import ols
import statsmodels.api as sm
import scipy.stats as stats
from sklearn.model_selection import train_test_split

#Create training and testing set
train, test = train_test_split(kcdf)
```

```
In [50]: #Check the train and testing sets
print(len(train), len(test))
```

```
14103 4701
```

```
In [51]: #Defining the problem
outcome = 'price'
x_cols = list(kcdf.columns)
x_cols.remove('price')
```

```
In [52]: #Fitting the model
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula = formula, data = train).fit()
model.summary()
```

Out[52]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.726
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.725
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1036.
<b>Date:</b>	Tue, 22 Dec 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	12:07:00	<b>Log-Likelihood:</b>	-1.9167e+05
<b>No. Observations:</b>	14103	<b>AIC:</b>	3.834e+05
<b>Df Residuals:</b>	14066	<b>BIC:</b>	3.837e+05
<b>Df Model:</b>	36		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	1.565e+06	2.03e+05	7.729	0.000	1.17e+06	1.96e+06
<b>bedrooms</b>	-2.046e+04	2242.097	-9.126	0.000	-2.49e+04	-1.61e+04
<b>bathrooms</b>	1.919e+04	2892.165	6.636	0.000	1.35e+04	2.49e+04
<b>sqftliving</b>	7.071e+04	1999.173	35.368	0.000	6.68e+04	7.46e+04
<b>sqftlot</b>	9520.0360	2385.760	3.990	0.000	4843.630	1.42e+04
<b>sqftabove</b>	5.758e+04	2175.561	26.469	0.000	5.33e+04	6.18e+04
<b>sqftbasement</b>	3.915e+04	1848.419	21.182	0.000	3.55e+04	4.28e+04
<b>lat</b>	8.07e+04	1754.975	45.981	0.000	7.73e+04	8.41e+04
<b>long</b>	-3.129e+04	1946.671	-16.071	0.000	-3.51e+04	-2.75e+04
<b>sqftliving15</b>	3.01e+04	2803.615	10.737	0.000	2.46e+04	3.56e+04
<b>sqftlot15</b>	-1.349e+04	2419.915	-5.575	0.000	-1.82e+04	-8746.860
<b>floors15</b>	5.755e+04	6230.414	9.236	0.000	4.53e+04	6.98e+04
<b>floors2</b>	-1.151e+04	5196.402	-2.215	0.027	-2.17e+04	-1325.757
<b>floors25</b>	1.463e+05	1.92e+04	7.615	0.000	1.09e+05	1.84e+05
<b>floors3</b>	-4697.1980	1.11e+04	-0.422	0.673	-2.65e+04	1.71e+04
<b>floors35</b>	41.8800	9.7e+04	0.000	1.000	-1.9e+05	1.9e+05
<b>waterfront1</b>	7.838e+05	1.9e+04	41.186	0.000	7.47e+05	8.21e+05
<b>condition2</b>	5.396e+04	4.96e+04	1.088	0.277	-4.32e+04	1.51e+05
<b>condition3</b>	5.405e+04	4.61e+04	1.173	0.241	-3.63e+04	1.44e+05
<b>condition4</b>	1.066e+05	4.61e+04	2.311	0.021	1.62e+04	1.97e+05
<b>condition5</b>	1.546e+05	4.64e+04	3.333	0.001	6.37e+04	2.45e+05
<b>grade4</b>	-1.614e+05	1.98e+05	-0.813	0.416	-5.5e+05	2.28e+05
<b>grade5</b>	-2.19e+05	1.94e+05	-1.127	0.260	-6e+05	1.62e+05
<b>grade6</b>	-2.217e+05	1.94e+05	-1.144	0.253	-6.01e+05	1.58e+05
<b>grade7</b>	-2.091e+05	1.94e+05	-1.079	0.281	-5.89e+05	1.71e+05

<b>grade8</b>	-1.647e+05	1.94e+05	-0.850	0.396	-5.45e+05	2.15e+05
<b>grade9</b>	-4.494e+04	1.94e+05	-0.232	0.817	-4.25e+05	3.35e+05
<b>grade1</b>	1.102e+05	1.94e+05	0.568	0.570	-2.7e+05	4.91e+05
<b>grade11</b>	3.453e+05	1.94e+05	1.775	0.076	-3.59e+04	7.27e+05
<b>grade12</b>	7.624e+05	1.96e+05	3.889	0.000	3.78e+05	1.15e+06
<b>grade13</b>	1.433e+06	2.05e+05	6.976	0.000	1.03e+06	1.84e+06
<b>publicsafety1</b>	1.299e+04	1.17e+04	1.112	0.266	-9910.973	3.59e+04
<b>governmentbuilding1</b>	-1.638e+04	4790.196	-3.420	0.001	-2.58e+04	-6993.779
<b>foodandrestaurants1</b>	-9.8e+05	3.45e+04	-28.390	0.000	-1.05e+06	-9.12e+05
<b>shoppingandentertainment1</b>	-5.447e+04	1.63e+04	-3.347	0.001	-8.64e+04	-2.26e+04
<b>hubsofttransport1</b>	9.275e+04	6429.997	14.425	0.000	8.01e+04	1.05e+05
<b>medical1</b>	-1.227e+04	3846.647	-3.190	0.001	-1.98e+04	-4730.965
<b>other1</b>	-3394.5891	5194.045	-0.654	0.513	-1.36e+04	6786.428
<b>Omnibus:</b>	9358.198	<b>Durbin-Watson:</b>	1.995			
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	435402.193			
<b>Skew:</b>	2.597	<b>Prob(JB):</b>	0.00			
<b>Kurtosis:</b>	29.720	<b>Cond. No.</b>	1.08e+16			

## Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 9.19e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

In [53]: *#The p-values of some of the variables are above the alpha of 0.05, therefore  
#these variables should be removed*

```
#Backward elimination
y = kcdf['price']
X = kcdf[x_cols]
pmax = 1
while (len(x_cols) > 0 ):
    p = []
    X_1 = X[x_cols]
    X_1 = sm.add_constant(X_1)
    model = sm.OLS(y, X_1).fit()
    p = pd.Series(model.pvalues.values[1:], index = x_cols)
    pmax = max(p)
    feature_with_p_max = p.idxmax()
    if(pmax > 0.05):
        x_cols.remove(feature_with_p_max)
    else:
        break
selected_features_BE = x_cols
print(selected_features_BE)
```

C:\Users\u-ana\Anaconda3\envs\learn-env\lib\site-packages\numpy\core\fromnumeric.py:2580: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead.

return ptp(axis=axis, out=out, \*\*kwargs)

```
['bedrooms', 'bathrooms', 'sqftliving', 'sqftlot', 'sqftabove', 'sqftbasement', 'lat', 'long', 'sqftliving15', 'sqftlot15', 'floors15', 'floors2', 'floors25', 'waterfront1', 'condition4', 'condition5', 'grade4', 'grade5', 'grade6', 'grade7', 'grade8', 'grade1', 'grade11', 'grade12', 'grade13', 'governmentbuilding1', 'foodandrestaurants1', 'shoppingandentertainment1', 'hubsofttransport1', 'medical1']
```

```
In [54]: #Re-testing the model after removing initial variables
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula = formula, data = train).fit()
model.summary()
```

Out[54]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.726
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.726
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1287.
<b>Date:</b>	Tue, 22 Dec 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	12:07:01	<b>Log-Likelihood:</b>	-1.9167e+05
<b>No. Observations:</b>	14103	<b>AIC:</b>	3.834e+05
<b>Df Residuals:</b>	14073	<b>BIC:</b>	3.836e+05
<b>Df Model:</b>	29		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	1.576e+06	3.74e+04	42.187	0.000	1.5e+06	1.65e+06
<b>bedrooms</b>	-2.042e+04	2235.915	-9.132	0.000	-2.48e+04	-1.6e+04
<b>bathrooms</b>	1.892e+04	2804.867	6.745	0.000	1.34e+04	2.44e+04
<b>sqftliving</b>	7.083e+04	1995.764	35.489	0.000	6.69e+04	7.47e+04
<b>sqftlot</b>	9438.2857	2383.759	3.959	0.000	4765.802	1.41e+04
<b>sqftabove</b>	5.761e+04	2166.574	26.592	0.000	5.34e+04	6.19e+04
<b>sqftbasement</b>	3.935e+04	1790.234	21.980	0.000	3.58e+04	4.29e+04
<b>lat</b>	8.038e+04	1737.519	46.262	0.000	7.7e+04	8.38e+04
<b>long</b>	-3.122e+04	1921.655	-16.248	0.000	-3.5e+04	-2.75e+04
<b>sqftliving15</b>	3.005e+04	2775.993	10.826	0.000	2.46e+04	3.55e+04
<b>sqftlot15</b>	-1.353e+04	2416.348	-5.600	0.000	-1.83e+04	-8794.300
<b>floors15</b>	5.782e+04	6183.716	9.350	0.000	4.57e+04	6.99e+04
<b>floors2</b>	-1.096e+04	4889.314	-2.241	0.025	-2.05e+04	-1374.130
<b>floors25</b>	1.469e+05	1.91e+04	7.688	0.000	1.09e+05	1.84e+05
<b>waterfront1</b>	7.838e+05	1.9e+04	41.195	0.000	7.46e+05	8.21e+05
<b>condition4</b>	5.27e+04	3963.946	13.294	0.000	4.49e+04	6.05e+04
<b>condition5</b>	1.008e+05	6330.739	15.922	0.000	8.84e+04	1.13e+05
<b>grade4</b>	-1.184e+05	4.41e+04	-2.681	0.007	-2.05e+05	-3.18e+04
<b>grade5</b>	-1.755e+05	1.77e+04	-9.936	0.000	-2.1e+05	-1.41e+05
<b>grade6</b>	-1.769e+05	9365.695	-18.891	0.000	-1.95e+05	-1.59e+05
<b>grade7</b>	-1.642e+05	6972.468	-23.543	0.000	-1.78e+05	-1.5e+05
<b>grade8</b>	-1.199e+05	6164.766	-19.450	0.000	-1.32e+05	-1.08e+05
<b>grade1</b>	1.55e+05	8845.966	17.521	0.000	1.38e+05	1.72e+05
<b>grade11</b>	3.897e+05	1.39e+04	28.029	0.000	3.62e+05	4.17e+05
<b>grade12</b>	8.068e+05	2.78e+04	29.063	0.000	7.52e+05	8.61e+05



grade13	1.479e+06	6.66e+04	22.194	0.000	1.35e+06	1.61e+06
governmentbuilding1	-1.513e+04	4650.889	-3.253	0.001	-2.42e+04	-6013.598
foodandrestaurants1	-9.838e+05	3.42e+04	-28.775	0.000	-1.05e+06	-9.17e+05
shoppingandentertainment1	-4.259e+04	1.26e+04	-3.387	0.001	-6.72e+04	-1.79e+04
hubsofttransport1	9.13e+04	5577.495	16.369	0.000	8.04e+04	1.02e+05
medical1	-1.182e+04	3817.424	-3.097	0.002	-1.93e+04	-4338.644
Omnibus:	9356.496	Durbin-Watson:	1.995			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	434867.838			
Skew:	2.596	Prob(JB):	0.00			
Kurtosis:	29.704	Cond. No.	1.12e+16			


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 6.29e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Check Multicollinearity Assumption

```
In [55]: #Checking for multicollinearity
plt.figure(figsize = [20, 10])
sns.heatmap(kc.corr(), annot = True).set_title('KCDF Colinearity Heatmap \n',
fontsize = 20, fontweight = 'bold');
```



```
In [56]: from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [57]: #Use the variance inflation factor to check for multicollinearity
X = kcdf[x_cols]
vif = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
list(zip(x_cols, vif))
```

C:\Users\u-ana\Anaconda3\envs\learn-env\lib\site-packages\statsmodels\stats\outliers\_influence.py:193: RuntimeWarning: divide by zero encountered in double\_scalars

```
vif = 1. / (1. - r_squared_i)
```

```
Out[57]: [('bedrooms', 1.7166227832652403),
('bathrooms', 2.9985776738318406),
('sqftliving', inf),
('sqftlot', 2.118808489483488),
('sqftabove', inf),
('sqftbasement', inf),
('lat', 1.1351959926540611),
('long', 1.3711379122874103),
('sqftliving15', 2.9017866221437223),
('sqftlot15', 2.1698167429792505),
('floors15', 1.2562869969373414),
('floors2', 3.4314395063081924),
('floors25', 1.0691784617657754),
('waterfront1', 1.0441222124176783),
('condition4', 1.5524025500582985),
('condition5', 1.1984172908282784),
('grade4', 1.0378914843955527),
('grade5', 1.274195467114227),
('grade6', 3.0903692457906584),
('grade7', 7.488973294056128),
('grade8', 3.9678980193278246),
('grade1', 1.5458974162235335),
('grade11', 1.356705416423),
('grade12', 1.1753421790812582),
('grade13', 1.0740266980649464),
('governmentbuilding1', 6.701529175565841),
('foodandrestaurants1', 73.66546234409783),
('shoppingandentertainment1', 52.57426767791397),
('hubsoftransport1', 10.020334789989056),
('medical1', 3.882635544418143)]
```

```
In [58]: #Remove highly colinear columns
cols_to_remove = ['sqftliving', 'sqftabove', 'sqftbasement', 'foodandrestaurant1', 'shoppingandentertainment1',
                  'hubsofttransport1']
for col in cols_to_remove:
    x_cols.remove(col)
print(x_cols)

['bedrooms', 'bathrooms', 'sqftlot', 'lat', 'long', 'sqftliving15', 'sqftlot15', 'floors15', 'floors2', 'floors25', 'waterfront1', 'condition4', 'condition5', 'grade4', 'grade5', 'grade6', 'grade7', 'grade8', 'grade1', 'grade11', 'grade12', 'grade13', 'governmentbuilding1', 'medical1']
```

```
In [59]: #Re-testing the model after addressing the colinearity issue
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula = formula, data = train).fit()
model.summary()
```

Out[59]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.676			
Model:	OLS	Adj. R-squared:	0.676			
Method:	Least Squares	F-statistic:	1226.			
Date:	Tue, 22 Dec 2020	Prob (F-statistic):	0.00			
Time:	12:07:03	Log-Likelihood:	-1.9285e+05			
No. Observations:	14103	AIC:	3.857e+05			
Df Residuals:	14078	BIC:	3.859e+05			
Df Model:	24					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.641e+05	7974.708	83.276	0.000	6.48e+05	6.8e+05
bedrooms	1.139e+04	2232.537	5.101	0.000	7011.098	1.58e+04
bathrooms	6.201e+04	2801.278	22.137	0.000	5.65e+04	6.75e+04
sqftlot	1.411e+04	2585.536	5.456	0.000	9039.868	1.92e+04
lat	8.703e+04	1852.804	46.970	0.000	8.34e+04	9.07e+04
long	-3.626e+04	2016.413	-17.983	0.000	-4.02e+04	-3.23e+04
sqftliving15	6.98e+04	2771.686	25.183	0.000	6.44e+04	7.52e+04
sqftlot15	-1.042e+04	2623.158	-3.971	0.000	-1.56e+04	-5273.654
floors15	6.595e+04	6592.522	10.003	0.000	5.3e+04	7.89e+04
floors2	-1.768e+04	4750.334	-3.722	0.000	-2.7e+04	-8368.937
floors25	1.677e+05	2.05e+04	8.160	0.000	1.27e+05	2.08e+05
waterfront1	8.337e+05	2.06e+04	40.472	0.000	7.93e+05	8.74e+05
condition4	6.828e+04	4287.447	15.925	0.000	5.99e+04	7.67e+04
condition5	1.234e+05	6832.983	18.059	0.000	1.1e+05	1.37e+05
grade4	-2.344e+05	4.78e+04	-4.901	0.000	-3.28e+05	-1.41e+05
grade5	-2.677e+05	1.9e+04	-14.120	0.000	-3.05e+05	-2.31e+05
grade6	-2.592e+05	9853.860	-26.304	0.000	-2.79e+05	-2.4e+05
grade7	-2.312e+05	7249.231	-31.899	0.000	-2.45e+05	-2.17e+05
grade8	-1.659e+05	6526.507	-25.422	0.000	-1.79e+05	-1.53e+05
grade1	2.065e+05	9454.307	21.845	0.000	1.88e+05	2.25e+05
grade11	5.274e+05	1.45e+04	36.380	0.000	4.99e+05	5.56e+05
grade12	1.057e+06	2.93e+04	36.022	0.000	9.99e+05	1.11e+06
grade13	1.988e+06	7.1e+04	27.987	0.000	1.85e+06	2.13e+06
governmentbuilding1	-2845.8584	5001.200	-0.569	0.569	-1.26e+04	6957.156
medical1	-1.4e+04	4115.148	-3.402	0.001	-2.21e+04	-5932.964

<b>Omnibus:</b>	10441.973	<b>Durbin-Watson:</b>	1.983
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	703850.588
<b>Skew:</b>	2.953	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	37.101	<b>Cond. No.</b>	67.5

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## Re-Checking the features of the model

```
In [60]: from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import RFE
```

```
In [61]: #Re-checking for feature elimination with RFE
model = LinearRegression()
```

```
#Initializing the RFE model
rfe = RFE(model)
```

```
#Transforming data using RFE
X = kcdf[x_cols]
X_rfe = rfe.fit_transform(X, y)
```

```
#Fitting the data to the model
model.fit(X_rfe, y)
```

```
print(x_cols)
print(rfe.support_)
print(rfe.ranking_)
```

```
['bedrooms', 'bathrooms', 'sqftlot', 'lat', 'long', 'sqftliving15', 'sqftlot1
5', 'floors15', 'floors2', 'floors25', 'waterfront1', 'condition4', 'conditio
n5', 'grade4', 'grade5', 'grade6', 'grade7', 'grade8', 'grade1', 'grade11',
'grade12', 'grade13', 'governmentbuilding1', 'medical1']
[False False False False False False False False False True True False
 True True True True True True True True True True False False]
[10  5 11  2  7  6 12  3  8  1  1  4  1  1  1  1  1  1  1  1 13  9]
```

```
In [62]: #Remove columns
cols_to_remove = ['bedrooms', 'bathrooms', 'sqftlot', 'lat', 'long', 'sqftliving', 'sqftlot15', 'floors15', 'floors2', 'condition4', 'governmentbuilding1', 'medical1']
for col in cols_to_remove:
    if col in x_cols:
        x_cols.remove(col)

print(x_cols)

['sqftliving15', 'floors25', 'waterfront1', 'condition5', 'grade4', 'grade5', 'grade6', 'grade7', 'grade8', 'grade1', 'grade11', 'grade12', 'grade13']
```

```
In [63]: #Re-testing the model
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula = formula, data = train).fit()
model.summary()
```



Out[63]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.583
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.583
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1516.
<b>Date:</b>	Tue, 22 Dec 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	12:07:04	<b>Log-Likelihood:</b>	-1.9463e+05
<b>No. Observations:</b>	14103	<b>AIC:</b>	3.893e+05
<b>Df Residuals:</b>	14089	<b>BIC:</b>	3.894e+05
<b>Df Model:</b>	13		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	7.031e+05	6394.776	109.946	0.000	6.91e+05	7.16e+05
<b>sqftliving15</b>	6.683e+04	2907.223	22.987	0.000	6.11e+04	7.25e+04
<b>floors25</b>	2.074e+05	2.3e+04	9.012	0.000	1.62e+05	2.53e+05
<b>waterfront1</b>	8.567e+05	2.32e+04	36.877	0.000	8.11e+05	9.02e+05
<b>condition5</b>	1.359e+05	7463.585	18.211	0.000	1.21e+05	1.51e+05
<b>grade4</b>	-4.33e+05	5.38e+04	-8.041	0.000	-5.38e+05	-3.27e+05
<b>grade5</b>	-4.27e+05	2.08e+04	-20.547	0.000	-4.68e+05	-3.86e+05
<b>grade6</b>	-3.675e+05	1.02e+04	-36.155	0.000	-3.87e+05	-3.48e+05
<b>grade7</b>	-2.86e+05	7688.398	-37.202	0.000	-3.01e+05	-2.71e+05
<b>grade8</b>	-1.842e+05	7312.056	-25.195	0.000	-1.99e+05	-1.7e+05
<b>grade1</b>	2.393e+05	1.07e+04	22.455	0.000	2.18e+05	2.6e+05
<b>grade11</b>	6.114e+05	1.62e+04	37.716	0.000	5.8e+05	6.43e+05
<b>grade12</b>	1.165e+06	3.3e+04	35.330	0.000	1.1e+06	1.23e+06
<b>grade13</b>	2.254e+06	8.01e+04	28.136	0.000	2.1e+06	2.41e+06

<b>Omnibus:</b>	9040.504	<b>Durbin-Watson:</b>	1.980
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	368023.165
<b>Skew:</b>	2.506	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	27.519	<b>Cond. No.</b>	45.8

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [64]: #Let's re-add some features
x_cols.append('lat')
```

```
In [65]: #Re-testing the model
predictors = '+' .join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula = formula, data = train).fit()
model.summary()
```

Out[65]: OLS Regression Results

Dep. Variable:	price		R-squared:		0.642	
Model:	OLS		Adj. R-squared:		0.642	
Method:	Least Squares		F-statistic:		1806.	
Date:	Tue, 22 Dec 2020		Prob (F-statistic):		0.00	
Time:	12:07:04		Log-Likelihood:		-1.9356e+05	
No. Observations:	14103		AIC:		3.871e+05	
Df Residuals:	14088		BIC:		3.873e+05	
Df Model:	14					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.876e+05	5933.212	115.898	0.000	6.76e+05	6.99e+05
sqftliving15	7.19e+04	2695.517	26.676	0.000	6.66e+04	7.72e+04
floors25	2.031e+05	2.13e+04	9.522	0.000	1.61e+05	2.45e+05
waterfront1	8.794e+05	2.15e+04	40.846	0.000	8.37e+05	9.22e+05
condition5	1.258e+05	6917.966	18.189	0.000	1.12e+05	1.39e+05
grade4	-3.621e+05	4.99e+04	-7.255	0.000	-4.6e+05	-2.64e+05
grade5	-3.668e+05	1.93e+04	-19.007	0.000	-4.05e+05	-3.29e+05
grade6	-3.261e+05	9457.316	-34.477	0.000	-3.45e+05	-3.08e+05
grade7	-2.619e+05	7140.673	-36.673	0.000	-2.76e+05	-2.48e+05
grade8	-1.73e+05	6778.394	-25.527	0.000	-1.86e+05	-1.6e+05
grade1	2.267e+05	9875.356	22.961	0.000	2.07e+05	2.46e+05
grade11	5.891e+05	1.5e+04	39.203	0.000	5.6e+05	6.19e+05
grade12	1.146e+06	3.06e+04	37.485	0.000	1.09e+06	1.21e+06
grade13	2.206e+06	7.42e+04	29.728	0.000	2.06e+06	2.35e+06
lat	9.052e+04	1876.897	48.229	0.000	8.68e+04	9.42e+04
Omnibus:	10404.395	Durbin-Watson:		1.992		
Prob(Omnibus):	0.000	Jarque-Bera (JB):		622079.394		
Skew:	2.977	Prob(JB):		0.00		
Kurtosis:	34.987	Cond. No.		45.8		


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## Checking for Outliers

```
In [66]: #Checking for outliers
plt.figure(figsize = (20, 5))
sns.boxplot(kc['price'])

plt.title('House Price Distribution Outliers \n', fontsize = 20, fontweight =
'bold');
```



```
In [67]: #Calculate Summary Statistics to eliminate outliers
price_mean, price_std = np.mean(kc.price), np.std(kc.price)

cut_off = price_std * 3
lower, upper = price_mean - cut_off, price_mean + cut_off
```

```
In [68]: #identify outliers
outliers = [x for x in kc.price if x < lower or x > upper]

print(len(outliers))
print(sorted(outliers)[:5])

344
[1670000.0, 1670000.0, 1680000.0, 1680000.0, 1680000.0]
```

```
In [69]: #Remove rows with outlier values
kcdf = kcdf[kcdf.price < 1670000]
```

In [70]: *#Check DataFrame information*

kcdf.info()

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 18460 entries, 0 to 18803
Data columns (total 38 columns):
price                                18460 non-null float64
bedrooms                            18460 non-null float64
bathrooms                           18460 non-null float64
sqftliving                           18460 non-null float64
sqftlot                              18460 non-null float64
sqftabove                           18460 non-null float64
sqftbasement                        18460 non-null float64
lat                                  18460 non-null float64
long                                 18460 non-null float64
sqftliving15                         18460 non-null float64
sqftlot15                           18460 non-null float64
floors15                            18460 non-null uint8
floors2                             18460 non-null uint8
floors25                            18460 non-null uint8
floors3                             18460 non-null uint8
floors35                            18460 non-null uint8
waterfront1                         18460 non-null uint8
condition2                          18460 non-null uint8
condition3                          18460 non-null uint8
condition4                          18460 non-null uint8
condition5                          18460 non-null uint8
grade4                              18460 non-null uint8
grade5                              18460 non-null uint8
grade6                              18460 non-null uint8
grade7                              18460 non-null uint8
grade8                              18460 non-null uint8
grade9                              18460 non-null uint8
grade1                              18460 non-null uint8
grade11                             18460 non-null uint8
grade12                             18460 non-null uint8
grade13                             18460 non-null uint8
publicsafety1                      18460 non-null uint8
governmentbuilding1                18460 non-null uint8
foodandrestaurants1                18460 non-null uint8
shoppingandentertainment1          18460 non-null uint8
hubsoftransport1                   18460 non-null uint8
medical1                           18460 non-null uint8
other1                              18460 non-null uint8
dtypes: float64(11), uint8(27)
memory usage: 2.2 MB

```

```
In [71]: #Re-testing the model  
#Create new training and testing set  
train, test = train_test_split(kcdf)  
  
#Fit  
predictors = '+' .join(x_cols)  
formula = outcome + '~' + predictors  
model = ols(formula = formula, data = train).fit()  
model.summary()
```

```
C:\Users\u-ana\Anaconda3\envs\learn-env\lib\site-packages\statsmodels\base\model.py:1362: RuntimeWarning: invalid value encountered in true_divide
  return self.params / self.bse
```

Out[71]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.624			
Model:	OLS	Adj. R-squared:	0.624			
Method:	Least Squares	F-statistic:	1766.			
Date:	Tue, 22 Dec 2020	Prob (F-statistic):	0.00			
Time:	12:07:04	Log-Likelihood:	-1.8565e+05			
No. Observations:	13845	AIC:	3.713e+05			
Df Residuals:	13831	BIC:	3.714e+05			
Df Model:	13					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.714e+05	4372.187	153.563	0.000	6.63e+05	6.8e+05
sqftliving15	6.54e+04	2047.908	31.934	0.000	6.14e+04	6.94e+04
floors25	1.339e+05	1.7e+04	7.862	0.000	1.01e+05	1.67e+05
waterfront1	4.09e+05	2.2e+04	18.567	0.000	3.66e+05	4.52e+05
condition5	1.172e+05	5161.453	22.705	0.000	1.07e+05	1.27e+05
grade4	-3.636e+05	3.56e+04	-10.218	0.000	-4.33e+05	-2.94e+05
grade5	-3.506e+05	1.44e+04	-24.325	0.000	-3.79e+05	-3.22e+05
grade6	-3.116e+05	7018.018	-44.394	0.000	-3.25e+05	-2.98e+05
grade7	-2.474e+05	5275.811	-46.893	0.000	-2.58e+05	-2.37e+05
grade8	-1.518e+05	4972.339	-30.521	0.000	-1.62e+05	-1.42e+05
grade1	1.36e+05	7540.894	18.040	0.000	1.21e+05	1.51e+05
grade11	3.093e+05	1.26e+04	24.614	0.000	2.85e+05	3.34e+05
grade12	4.564e+05	3.85e+04	11.846	0.000	3.81e+05	5.32e+05
grade13	0	0	nan	nan	0	0
lat	8.614e+04	1376.766	62.570	0.000	8.34e+04	8.88e+04
Omnibus:	3746.133	Durbin-Watson:	2.021			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	13329.345			
Skew:	1.336	Prob(JB):	0.00			
Kurtosis:	6.996	Cond. No.	2.63e+19			


Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.68e-35. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.



## Check Normality Assumption

```
In [72]: #Check for normality assumption with qqplot
residuals = model.resid
fig = sm.graphics.qqplot(residuals, dist = stats.norm, line = '45', fit = True
)
```



```
In [73]: #The right tail seems to still contain outlier values
#Drop outliers

#Find percentile cutoff point
for i in range(90, 100):
    q = i / 100
    print('{} percentile: {}'.format(q, kcdf['price'].quantile(q=q)))
```

```
0.9 percentile: 850000.0
0.91 percentile: 870000.0
0.92 percentile: 900000.0
0.93 percentile: 930000.0
0.94 percentile: 969229.9999999995
0.95 percentile: 1010000.0
0.96 percentile: 1100000.0
0.97 percentile: 1200000.0
0.98 percentile: 1300000.0
0.99 percentile: 1430000.0
```

```
In [74]: #Remove rows with outlier values from 98th percentile up
#kcdf = kcdf[kcdf.price < 1300000]
```

```
In [75]: #Re-testing the model  
#Create new training and testing set  
train, test = train_test_split(kcdf)  
  
#Fit  
predictors = '+' .join(x_cols)  
formula = outcome + '~' + predictors  
model = ols(formula = formula, data = train).fit()  
model.summary()
```

```
C:\Users\u-ana\Anaconda3\envs\learn-env\lib\site-packages\statsmodels\regression\linear_model.py:1827: RuntimeWarning: divide by zero encountered in double_scalars
  return np.sqrt(eigvals[0]/eigvals[-1])
```

Out[75]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.627
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.627
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1789.
<b>Date:</b>	Tue, 22 Dec 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	12:07:05	<b>Log-Likelihood:</b>	-1.8555e+05
<b>No. Observations:</b>	13845	<b>AIC:</b>	3.711e+05
<b>Df Residuals:</b>	13831	<b>BIC:</b>	3.712e+05
<b>Df Model:</b>	13		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	6.723e+05	4346.193	154.680	0.000	6.64e+05	6.81e+05
<b>sqftliving15</b>	6.693e+04	2026.022	33.034	0.000	6.3e+04	7.09e+04
<b>floors25</b>	1.342e+05	1.71e+04	7.851	0.000	1.01e+05	1.68e+05
<b>waterfront1</b>	4.087e+05	2.13e+04	19.200	0.000	3.67e+05	4.5e+05
<b>condition5</b>	1.124e+05	5103.373	22.018	0.000	1.02e+05	1.22e+05
<b>grade4</b>	-3.668e+05	3.31e+04	-11.088	0.000	-4.32e+05	-3.02e+05
<b>grade5</b>	-3.409e+05	1.39e+04	-24.519	0.000	-3.68e+05	-3.14e+05
<b>grade6</b>	-3.12e+05	6878.522	-45.357	0.000	-3.25e+05	-2.99e+05
<b>grade7</b>	-2.483e+05	5231.052	-47.458	0.000	-2.59e+05	-2.38e+05
<b>grade8</b>	-1.536e+05	4958.853	-30.973	0.000	-1.63e+05	-1.44e+05
<b>grade1</b>	1.33e+05	7570.333	17.573	0.000	1.18e+05	1.48e+05
<b>grade11</b>	2.904e+05	1.25e+04	23.153	0.000	2.66e+05	3.15e+05
<b>grade12</b>	4.833e+05	3.4e+04	14.218	0.000	4.17e+05	5.5e+05
<b>grade13</b>	0	0	nan	nan	0	0
<b>lat</b>	8.614e+04	1358.564	63.403	0.000	8.35e+04	8.88e+04


<b>Omnibus:</b>	3641.138	<b>Durbin-Watson:</b>	1.988
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	12682.470
<b>Skew:</b>	1.305	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	6.896	<b>Cond. No.</b>	inf

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 0. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [76]: plt.figure(figsize = [8, 6])

sns.distplot(kcdf['price'], color = 'c').set_title('Housing Prices Distributio
n', fontsize = 20)
plt.xlabel('\n Price in USD', fontsize = 12)
plt.ylabel('Frequency \n', fontsize = 12);
```



Removing some rows of data has decreased the  $r\_squared$ , therefore we revert to the previous model  $R\_2 = .666$

In [77]: *#The p-values of some of the variables are above the alpha of 0.05, therefore  
#we tried removing these variable*

```
#Backward elimination
y = kcdf['price']
X = kcdf[x_cols]
pmax = 1
while (len(x_cols) > 0 ):
    p = []
    X_1 = X[x_cols]
    X_1 = sm.add_constant(X_1)
    model = sm.OLS(y, X_1).fit()
    p = pd.Series(model.pvalues.values[1:], index = x_cols)
    pmax = max(p)
    feature_with_p_max = p.idxmax()
    if(pmax > 0.05):
        x_cols.remove(feature_with_p_max)
    else:
        break
selected_features_BE = x_cols
print(selected_features_BE)
```

```
['sqftliving15', 'floors25', 'waterfront1', 'condition5', 'grade4', 'grade5',
'grade6', 'grade7', 'grade8', 'grade1', 'grade11', 'grade12', 'grade13', 'lat']
```

```
C:\Users\u-ana\Anaconda3\envs\learn-env\lib\site-packages\numpy\core\fromnumeric.py:2580: FutureWarning: Method .ptp is deprecated and will be removed in
a future version. Use numpy.ptp instead.
    return ptp(axis=axis, out=out, **kwargs)
```

```
In [78]: #Re-Check model  
predictors = '+'.join(x_cols)  
formula = outcome + '~' + predictors  
model = ols(formula = formula, data = train).fit()  
model.summary()
```

Out[78]: OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.627
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.627
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1789.
<b>Date:</b>	Tue, 22 Dec 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	12:07:05	<b>Log-Likelihood:</b>	-1.8555e+05
<b>No. Observations:</b>	13845	<b>AIC:</b>	3.711e+05
<b>Df Residuals:</b>	13831	<b>BIC:</b>	3.712e+05
<b>Df Model:</b>	13		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	6.723e+05	4346.193	154.680	0.000	6.64e+05	6.81e+05
<b>sqftliving15</b>	6.693e+04	2026.022	33.034	0.000	6.3e+04	7.09e+04
<b>floors25</b>	1.342e+05	1.71e+04	7.851	0.000	1.01e+05	1.68e+05
<b>waterfront1</b>	4.087e+05	2.13e+04	19.200	0.000	3.67e+05	4.5e+05
<b>condition5</b>	1.124e+05	5103.373	22.018	0.000	1.02e+05	1.22e+05
<b>grade4</b>	-3.668e+05	3.31e+04	-11.088	0.000	-4.32e+05	-3.02e+05
<b>grade5</b>	-3.409e+05	1.39e+04	-24.519	0.000	-3.68e+05	-3.14e+05
<b>grade6</b>	-3.12e+05	6878.522	-45.357	0.000	-3.25e+05	-2.99e+05
<b>grade7</b>	-2.483e+05	5231.052	-47.458	0.000	-2.59e+05	-2.38e+05
<b>grade8</b>	-1.536e+05	4958.853	-30.973	0.000	-1.63e+05	-1.44e+05
<b>grade1</b>	1.33e+05	7570.333	17.573	0.000	1.18e+05	1.48e+05
<b>grade11</b>	2.904e+05	1.25e+04	23.153	0.000	2.66e+05	3.15e+05
<b>grade12</b>	4.833e+05	3.4e+04	14.218	0.000	4.17e+05	5.5e+05
<b>grade13</b>	0	0	nan	nan	0	0
<b>lat</b>	8.614e+04	1358.564	63.403	0.000	8.35e+04	8.88e+04

<b>Omnibus:</b>	3641.138	<b>Durbin-Watson:</b>	1.988
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	12682.470
<b>Skew:</b>	1.305	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	6.896	<b>Cond. No.</b>	inf

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 0. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.



```
In [79]: x_cols.remove('grade13')
```

```
In [80]: #Re-Check model  
predictors = '+'.join(x_cols)  
formula = outcome + '~' + predictors  
model = ols(formula = formula, data = train).fit()  
model.summary()
```

Out[80]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.627			
Model:	OLS	Adj. R-squared:	0.627			
Method:	Least Squares	F-statistic:	1789.			
Date:	Tue, 22 Dec 2020	Prob (F-statistic):	0.00			
Time:	12:07:05	Log-Likelihood:	-1.8555e+05			
No. Observations:	13845	AIC:	3.711e+05			
Df Residuals:	13831	BIC:	3.712e+05			
Df Model:	13					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	6.723e+05	4346.193	154.680	0.000	6.64e+05	6.81e+05
sqftliving15	6.693e+04	2026.022	33.034	0.000	6.3e+04	7.09e+04
floors25	1.342e+05	1.71e+04	7.851	0.000	1.01e+05	1.68e+05
waterfront1	4.087e+05	2.13e+04	19.200	0.000	3.67e+05	4.5e+05
condition5	1.124e+05	5103.373	22.018	0.000	1.02e+05	1.22e+05
grade4	-3.668e+05	3.31e+04	-11.088	0.000	-4.32e+05	-3.02e+05
grade5	-3.409e+05	1.39e+04	-24.519	0.000	-3.68e+05	-3.14e+05
grade6	-3.12e+05	6878.522	-45.357	0.000	-3.25e+05	-2.99e+05
grade7	-2.483e+05	5231.052	-47.458	0.000	-2.59e+05	-2.38e+05
grade8	-1.536e+05	4958.853	-30.973	0.000	-1.63e+05	-1.44e+05
grade1	1.33e+05	7570.333	17.573	0.000	1.18e+05	1.48e+05
grade11	2.904e+05	1.25e+04	23.153	0.000	2.66e+05	3.15e+05
grade12	4.833e+05	3.4e+04	14.218	0.000	4.17e+05	5.5e+05
lat	8.614e+04	1358.564	63.403	0.000	8.35e+04	8.88e+04
Omnibus:	3641.138	Durbin-Watson:	1.988			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	12682.470			
Skew:	1.305	Prob(JB):	0.00			
Kurtosis:	6.896	Cond. No.	29.0			

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The removal of outliers have reduced the R-squared values, therefore we will return to the prior model as the final one at present